

# Dynamic Programming

Weizi Li



University of Tennessee, Knoxville

- Introduction
- Policy Evaluation
- Policy Iteration
- Value Iteration
- Extensions of DP

# Introduction

- Reinforcement Learning: agent improves its policy by **interacting** with an **initially unknown** environment
  - ▶ Model-free learning: agent doesn't infer the model of the environment during learning
  - ▶ Model-based learning: agent infers the model of the environment and uses the inferred model during learning

- Planning: agent improves its policy by **computing** with its **internal model** of the environment (model-based)



- Dynamic Programming (DP) assumes full knowledge of the MDP (including the model, i.e., transition function and reward function) and is used for **planning** in the MDP
- DP's focus is to solve MDP using the least computational effort

- DP is a general method for problems that have two properties:
  - ▶ 1) Optimal substructure: optimal solution can be decomposed for subproblems
  - ▶ 2) Overlapping subproblems: subproblems recur many times and their solutions can be reused
- DP solves a problem by breaking it down into subproblems
  - ▶ solve the subproblems
  - ▶ combine solutions to subproblems
- Example problem: finding the shortest path between A and B

- Scheduling algorithms
- String algorithms (e.g., sequence alignment)
- Graph algorithms (e.g., shortest path algorithms)
- Graphical models (e.g., Viterbi algorithm)
- Bioinformatics (e.g., lattice models)
- ...




- MDPs have both properties required by DP:
  - ▶ Bellman equation gives recursive decomposition
  - ▶ Value function stores solutions for reuse
- An optimal policy can be subdivided into two components:
  - ▶ An optimal first action  $A^*$
  - ▶ Followed by an optimal policy from successor state  $S'$

- Prediction: evaluate a given policy (evaluate the future)
  - ▶ E.g., what is the value function for the uniform random policy?
  - ▶ Input: MDP and policy  $\pi$
  - ▶ Output: value function  $V^\pi$
- Control: find the best policy (optimize the future)
  - ▶ What is the optimal value function over all possible policies (hence the optimal policy)?
  - ▶ Input: MDP
  - ▶ Output: optimal value function  $V^*$  and optimal policy  $\pi^*$

# Policy Evaluation

- Policy evaluation solves the **prediction** problem: compute the state-value function  $V^\pi$  for an arbitrary policy  $\pi$
- Recall: one **direct** method to solve Bellman expectation equation is via the normal equation (only works for small MDPs)
- The DP approach is to **iteratively** applying Bellman expectation equation until converge to  $V^\pi$

- Initialize  $V_0^\pi(s) = 0$
- At the  $(k+1)$ th iteration, update  $V_{k+1}^\pi(s)$  from  $V_k^\pi(s')$
- DP uses *bootstrapping*: updates a guess ( $V_{k+1}^\pi(s)$ ) using a guess ( $V_k^\pi(s')$ )



$$v_\pi(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left( \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_\pi(s') \right)$$

- Evaluate the uniform random policy in an undiscounted, episodic task
  - ▶ Actions that would take the agent off the grid will leave the state unchanged

**Example 4.1** Consider the  $4 \times 4$  gridworld shown below.



	1	2	3
4	5	6	7
8	9	10	11
12	13	14	

$R_t = -1$   
on all transitions

- {Row 3, Column 4} from  $k = 2$  to  $k = 3$
- $-1.7 \rightarrow -2.4$ :  
 $-0.25(0+2+2+1.7)-1 = -2.425$


	$v_k$ for the random policy			
$k = 0$	0.0	0.0	0.0	0.0
	0.0	0.0	0.0	0.0
	0.0	0.0	0.0	0.0
	0.0	0.0	0.0	0.0
$k = 1$	0.0	-1.0	-1.0	-1.0
	-1.0	-1.0	-1.0	-1.0
	-1.0	-1.0	-1.0	-1.0
	-1.0	-1.0	-1.0	0.0
$k = 2$	0.0	-1.7	-2.0	-2.0
	-1.7	-2.0	-2.0	-2.0
	-2.0	-2.0	-2.0	-1.7
	-2.0	-2.0	-1.7	0.0
$k = 3$	0.0	-2.4	-2.9	-3.0
	-2.4	-2.9	-3.0	-2.9
	-2.9	-3.0	-2.9	-2.4
	-3.0	-2.9	-2.4	0.0
$k = 10$	0.0	-6.1	-8.4	-9.0
	-6.1	-7.7	-8.4	-8.4
	-8.4	-8.4	-7.7	-6.1
	-9.0	-8.4	-6.1	0.0
$k = \infty$	0.0	-14.	-20.	-22.
	-14.	-18.	-20.	-20.
	-20.	-20.	-18.	-14.
	-22.	-20.	-14.	0.0

- In theory, iterative policy evaluation converges in the limit:  
 $V_k^\pi \rightarrow V^\pi$  as  $k \rightarrow \infty$
- In practice, we stop when  $\|V_{k+1}^\pi - V_k^\pi\| \leq \epsilon$
- The existence and uniqueness of  $V^\pi$  are guaranteed providing:
  - ▶  $\gamma < 1$ ; or
  - ▶ any trajectory from any state under  $\pi$  is finite




# Policy Iteration

- Policy iteration solves the **control** problem, i.e., finding  $\pi^*$ 
  - ▶ policy evaluation: compute  $V^\pi$  for a given  $\pi$
  - ▶ policy improvement: produce  $\pi' \geq \pi$  by acting greedily to  $V^\pi$
- Policy iteration will converge to  $\pi^*$



- For a simple task, we can get  $\pi^*$  with one iteration of policy evaluation and policy improvement
- For the example in the figure, we can obtain  $\pi^*$  without converging to  $V^\pi$  in the first iteration of policy evaluation



- One drawback of the vanilla policy iteration is that each iteration involves policy evaluation, which can be computationally expensive if the number of states is large

- Do we need to converge to  $V^\pi$  to obtain the optimal policy?  
No, we don't
  - ▶ Use a stopping condition, e.g.,  $\|V_{k+1}^\pi - V_k^\pi\| \leq \epsilon$
  - ▶ Simply stop after  $k$  iterations (we may then arrive at the optimal policy)
- In theory, the policy evaluation step can be truncated in several ways without losing the convergence guarantees of policy iteration


- Do we need to update policy in every iteration to obtain the optimal policy? No, we can use value iteration

# Value Iteration

- Value iteration solves the **control** problem, i.e., finding  $\pi^*$
- Solution: iteratively applying Bellman optimality equation until converge to  $V^*$
- No direct method exists due to the nonlinearity of Bellman optimality equation
- Unlike policy iteration, there is no explicit policy update during the process
- Intermediate value functions may not correspond to any effective policy



- Initialize  $V_0^\pi(s) = 0$
- At the  $(k+1)$ th iteration, update  $V_{k+1}(s)$  from  $V_k(s')$ .



- Problem: find the shortest path from any state to the goal state (g)
  - ▶ Agent will stay if actions would take the agent off the grid
  - ▶  $R_t = -1$  on all transitions

g			

Problem

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

 $V_1$ 

0	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1

 $V_2$ 

0	-1	-2	-2
-1	-2	-2	-2
-2	-2	-2	-2
-2	-2	-2	-2

 $V_3$ 

0	-1	-2	-3
-1	-2	-3	-3
-2	-3	-3	-3
-3	-3	-3	-3

 $V_4$ 

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-4
-3	-4	-4	-4

 $V_5$ 

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-5
-3	-4	-5	-5

 $V_6$ 

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-5
-3	-4	-5	-6

 $V_7$

## ■ Policy iteration

- ▶ policy evaluation (multiple iterations) + policy improvement (multiple iterations)
- ▶ uses Bellman expectation equation

## ■ Value iteration

- ▶ finding optimal value function (multiple iterations) + optimal policy extraction (once)
- ▶ uses Bellman optimality equation

## ■ The performance is task-dependent

- ▶ policy converge rate (policy iteration) vs. value convergence rate (value iteration)
- ▶ if optimal policy is unnecessary, we can use policy iteration

- Both are widely used and converge much faster than their theoretical worst-case run times
- There is a spectrum of algorithms between Policy Iteration and Value Iteration conducting policy evaluation and policy improvement at different granularity
- In theory, initialization is independent of convergence rate; in practice, good initial value functions or policies can speed-up the convergence

# Extensions of DP

- Synchronous DP processes all states at once using previous value functions.

Problem	Bellman Equation	Algorithm
Prediction	Bellman Expectation Equation	Iterative Policy Evaluation
Control	Bellman Expectation Equation + Greedy Policy Improvement	Policy Iteration
Control	Bellman Optimality Equation	Value Iteration

- For  $m$  actions and  $n$  states:
  - ▶ Algorithms based on state-value function  $V$ , the complexity is  $O(mn^2)$  per iteration
  - ▶ Algorithms based on action-value function  $Q$ , the complexity is  $O(m^2n^2)$  per iteration
- DP is guaranteed to find an optimal policy in polynomial time, even though the total number of (deterministic) policies is  $m^n$

- Asynchronous DP does not require processing all states at once, instead it can select any state to process, in any order
- Asynchronous DP can significantly reduce the computation and guarantee to converge if all states continue to be selected
- Example asynchronous DP algorithms:
  - ▶ In-place dynamic programming
  - ▶ Prioritized sweeping
  - ▶ Real-time dynamic programming



- DP uses full-width updates by considering every successor state and action
- DP is effective and efficient for medium-sized problems (millions of states)
- For large problems, DP suffers Bellman's curse of dimensionality (the number of states grows exponentially with the number of state variables); in this case, even one update can be cost prohibitive

- Update via sampling the environment transition function and reward function has two advantages
  - ▶ Model-free: no need to know the full MDP
  - ▶ Overcomes the curse of dimensionality: cost is constant, independent from the number of state variables